

Backpropagation

Intuitive Explanation

- **Backpropagation** is a method for computing gradients
 - *Calculates how a single example would want to change the weights and biases*
 - How much and in which direction
 - **Gradient Descent** - carrying out backpropagation on all training data, averaging out gradients, each epoch
 - **SGD** - do backpropagation on each piece of **minibatch**, averaging fewer gradients, **faster** epochs
- Improve performance by considering changes to the **bias**, the **weights** and the **activation** of the previous layer
- We start from our output and calculate the weights backwards
- We want to most change the neurons that affect our output the most, i.e. the gradient of which is largest, i.e. are the most wrong
- At each node, we add all of the gradients / changes coming from the nodes in the next layer
 - I.e. at each node of the final hidden layer, we add up all of the gradients which are the opinions of each output node on how the output should change
 - This gives us the nudge that we want the nodes in the previous layer to result in

$$\delta^L = \nabla_a C \odot \sigma'(z^L)$$

$$\delta^l = ((w^{l+1})^T \delta^{l+1}) \odot \sigma'(z^l)$$

$$\frac{\partial C}{\partial b_j^l} = \delta_j^l$$

$$\frac{\partial C}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l$$

Backpropagation Calculus

- $z^L = W^T X + B$, the input into activation function at a neuron in layer L
- L - output layer
- a^L - the output of the activation function, $\sigma(z^L)$
 - a_1^L - the activation of the first neuron in the output layer
- C_0 - the cost of 0th training example, $(a^L - y)^2$
 - $C = \text{avg}(\sum_k (dC_k / dw^L))$
 - **∇ Nabla (gradient vector) - a vector of partial derivatives for f**
 - ∇C - a vector $[\partial C / \partial w^1, \partial C / \partial b^1, \dots, \partial C / \partial w^L, \partial C / \partial b^L]$
 - Tells us the gradient of every single parameter

- **Backpropagation** is finding out what is the gradient of the cost function with respect to the weights connected to the nodes in the output layer

- We use chain rule to look to go from the derivative of C_0 wrt a^L to the derivative of z^L wrt w^L
- $dz^L / dw^L = a^{L-1}$ tells us that the influence of w on z depends on previous layer's activation a^{L-1} , which makes sense
- To get the gradients for B , we substitute dz^L / dw^L for dz^L / db^L
- To propagate errors backwards, we get the gradient at a^{L-1} in a similar way, then we find the gradients of a^{L-1} wrt w^{L-1}
- We compute one layer of the gradients at a time
-

$$\frac{\partial C_0}{\partial a^{(L)}} = 2(a^{(L)} - y)$$

$$\frac{\partial a^{(L)}}{\partial z^{(L)}} = \sigma'(z^{(L)})$$

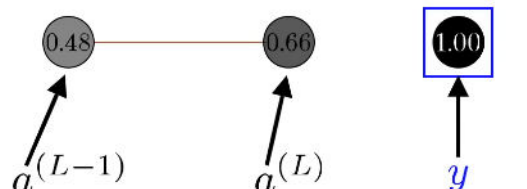
$$\frac{\partial z^{(L)}}{\partial w^{(L)}} = a^{(L-1)}$$

Equations for a neural network with one node in each layer

$$\text{Cost} \rightarrow C_0(\dots) = (a^{(L)} - y)^2$$

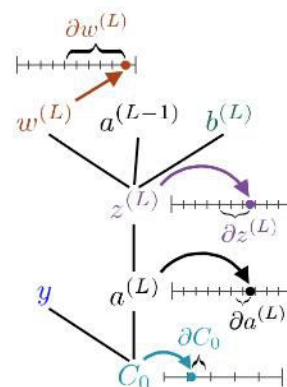
$$z^{(L)} = w^{(L)} a^{(L-1)} + b^{(L)}$$

$$a^{(L)} = \sigma(z^{(L)})$$



$$\frac{\partial C_0}{\partial w^{(L)}} = \frac{\partial z^{(L)}}{\partial w^{(L)}} \frac{\partial a^{(L)}}{\partial z^{(L)}} \frac{\partial C_0}{\partial a^{(L)}}$$

Chain rule



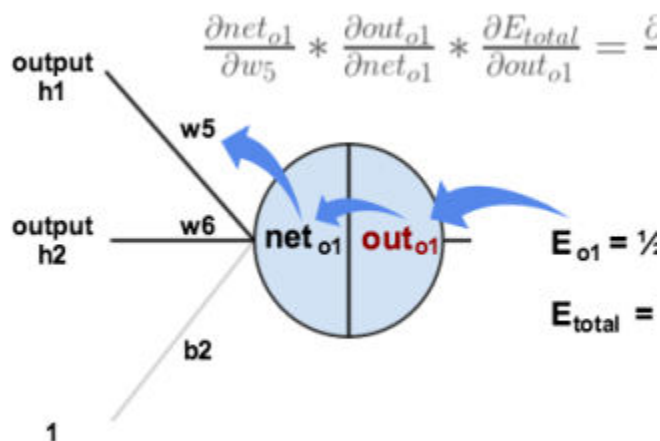
Equations for a normal neural network

- We must sum over the numerous neurons we have now
- dC_0/da_k^{L-1} becomes the sum of gradients over all nodes in L
- C_0 becomes the sum of errors across all output nodes in L

$$\frac{\partial C_0}{\partial w_{jk}^{(L)}} = \frac{\partial z_j^{(L)}}{\partial w_{jk}^{(L)}} \frac{\partial a_j^{(L)}}{\partial z_j^{(L)}} \frac{\partial C_0}{\partial a_j^{(L)}}$$

$$C_0 = \sum_{j=0}^{n_L-1} (a_j^{(L)} - y_j)^2$$

$$\frac{\partial C_0}{\partial a_k^{(L-1)}} = \underbrace{\sum_{j=0}^{n_L-1} \frac{\partial z_j^{(L)}}{\partial a_k^{(L-1)}} \frac{\partial a_j^{(L)}}{\partial z_j^{(L)}} \frac{\partial C_0}{\partial a_j^{(L)}}}_{\text{Sum over layer L}}$$



Extra

Parametric vs Nonparametric

- **Parametric ML Algorithms** - Number of parameters is fixed w.r.t the data sample size
 - Doesn't change its mind about how many parameters it will model
 - Complexity is bound by the number of parameters even if the amount of data is unbounded
 - All of the knowledge the model will learn is encapsulated by the **finite set of parameters**, which means that it is **independent of previously observed data**

$$P(x|\theta, \mathcal{D}) = P(x|\theta)$$

- E.g. In **OLS Regression**, number of parameters will always be length of $\beta + 1$ (variance)
 - E.g. simple linear regression
 - E.g. neural net with **fixed architecture** and **no weight decay**
 - *Greatly simplifies learning process, but limits what can be learned*
 - *Simple functions only*
 - E.g. linear regression, logistic regression, LDA, perceptrons
- **Nonparametric ML Algorithms** - The number of parameters can grow with the sample size
 - Assumes the data distribution cannot be defined in terms of a finite set of parameters
 - *Good when you have lots of data with no prior knowledge*
- **Argument they're the same** - a nonparametric algorithm can be approximated with a parametric model with an infinite number of parameters