

Introduction

- Tries to get rid of the necessary **negative samples** when doing **contrastive loss** for self-supervised learning
- They combine **momentum contrast** and ____ and then remove **negative samples**

Image Representation Learning - taking an image and feeding it through a function (e.g. NN like RN50) and make it give you a **representation vector h** that can be used to **solve many tasks**

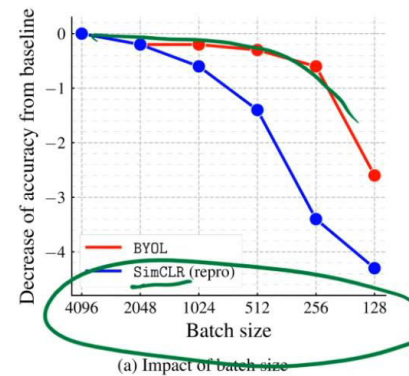
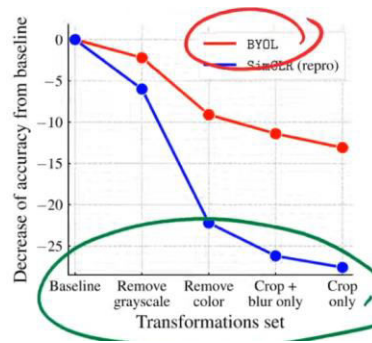
- Fine tune the model to give you good performance
- It means you can learn on a lot of data for something, then finetune to something where you have less data
- Similar approach to the ones of transformers like BERT

Self-Supervision - you'll take an image and make variants of that image using **augmentation**

- You exploit the fact that the different variants are the **same image** which should have **similar representations**

BYOL - has two neural networks, the **online network** and the **target network**

- Each of the two networks gets a **different augmentation** of the same image and the **online network** tries to **predict the target network representation**
- **Once you do it one way, you then do it with the same images but insert them in the opposite networks**
- The **target network** is trained with an **exponentially-moving average** of the online network
- **When it comes to blurring, i.e. gaussian blurring, we always blur the online network, but $p=0.1$ for target**
- The batch size is 4096 which requires **tons of computation**
 - Decreasing the batch size reduces accuracy but not as much as solutions with **negative samples**
 - **Negative samples are taken from the minibatch**
 - If you have less samples in your minibatch, you have less samples as a representation of the distribution your dataset so you have and you have **less negative samples** = worse performance
- More robust for **removal of certain augmentations**
- As **augmentations are important**, if you want to use other **modalities** (audio, text, video, ...) you need to find **similarly useful augmentations**



Problem - usually if you'd do this, you would get **collapse**


- The networks would cheat and make a constant function, e.g. $h = 0$ to always have the perfect loss

Old Solutions


- **Negative Sampling** - you take a different image and create a **third input** by augmenting it
 - The task now becomes to make the **first two** inputs have a **similar h** that's **different to the third**
 - The network can't map everything to the same function anymore
 - **Used as a combination with augmentations**
 - The negative sampling prevents **collapse**, but still allows the **augmentations** to help create **good representations**
 - **The augmentation choice is very important**

Idea - get rid of negative samples

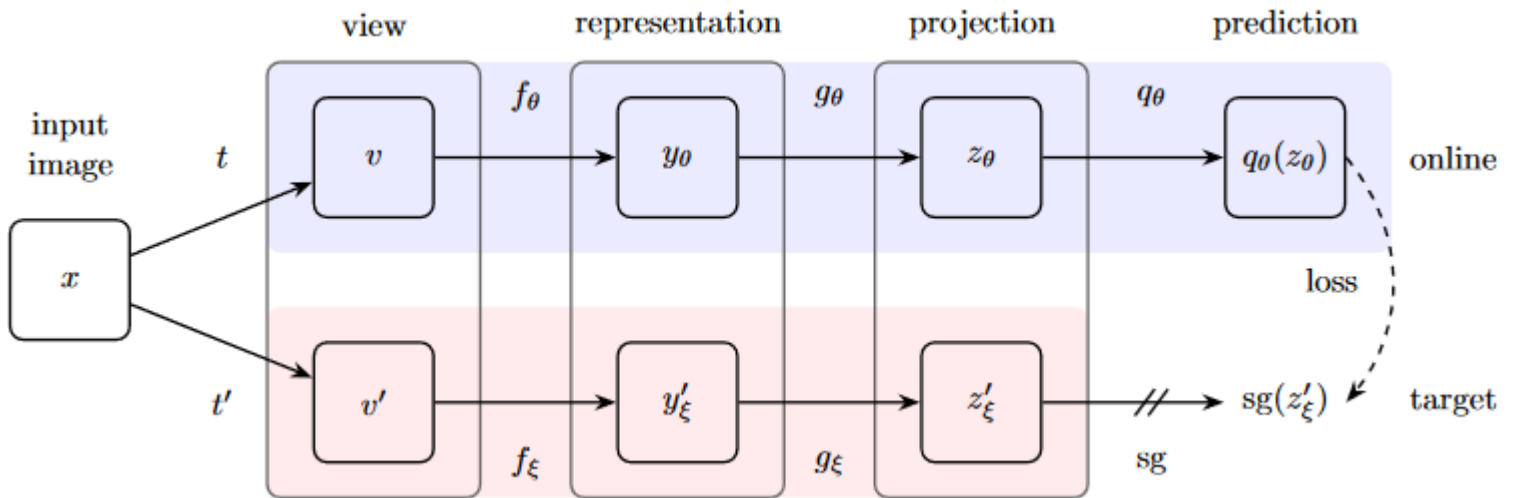
- There are many **issues** with **negative samples**
- Where do we get them? How do we sample them?
- This means that nothing is preventing $h = \text{constant } c$
- *Is this a super delicate balance?*



Many hyperparameters



- Is it just that you start in a local minimum?



BYOL Solution

- t and t' are two different **random augmentations** that create different **variants** v and v' that should be close
 - The **augmentations are responsible** for creating variants that are close but different and are responsible for creating really good **h representations**
- The online and target networks have the **same encoder with different parameters**
- We only **learn the online** parameters θ , the **target** parameters are a **copy** of the online parameters
 - **Momentum contrast** principle - Target parameters are a **lagging average** of the online parameters
 - You need to have a **stable representation of the past for a target**
 - We don't quite know why this works so well
- y and y' are two different **representations** created with our **encoders**
- f_θ is our **representer**, the output of this process
 - We discard everything else
- **Projections** are done to **reduce dimensionality** by making the **representation smaller**
 - We could **get rid of it if we wanted to**, there is no reason it's there apart from that it works
 - The first FF layer pumps up the dimensionality (say from 2048 \rightarrow 4096), the second FF goes down to 256
- We put the **representations** z , z' through the **predictor** q_θ which tries to take one input and predict the representation of the other input
 - We want $q_\theta(z) = z'$
 - Expanding that, we get $q_\theta(g_\theta(f_\theta(t(x)))) = g_\xi(f_\xi(t'(x'))))$
 - q_θ tries to nullify the difference in augmentations and representers
 - **q tries to predict the average $f_\xi(t(x))$ for all t** (independent of t)
- This means the **augmentations destroy all non-semantic information**
 - Therefore, our representations should contain only semantic info

et

$$q_\theta(f_\theta(A(x))) = f_\xi(A'(x))$$

$$E_A[f(A(x))]$$

$$\text{BYOL}_\theta \triangleq \|\overline{q_\theta(z_\theta)} - \overline{z'_\xi}\|_2^2 = 2 - 2 \cdot \frac{\langle q_\theta(z_\theta), z'_\xi \rangle}{\|q_\theta(z_\theta)\|_2 \cdot \|z'_\xi\|_2}$$

- You want the **L2 norm** of the z' representation to be close to the norm of the $q_\theta(z)$ **L2 representation**