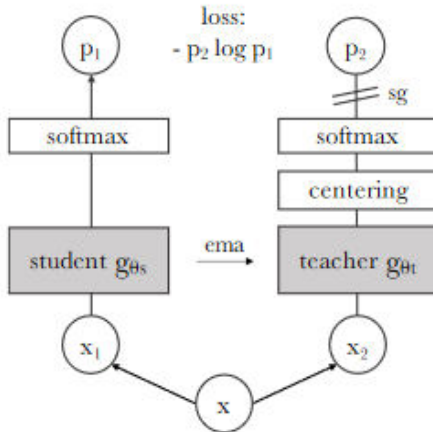


**DINO** - a method for unsupervised pre-training of ViTs

- **Self-Distillation with no labels**
- The transformer doesn't get taught what a class is, nor is it taught any form of image segmentation
- Can track objects behind occlusion
- If you use DINO for image classification and you view the **feature representation**, it will **cluster images of the same class** together, and it will also **cluster clusters of similar classes together**
  - **Feature Space -**
- Useful for things like *image retrieval (similar images clustered together)*, *classification*, *zero-shot classifier (by doing knn classifier)*, etc
- It does all of this with no supervision



### Architecture

- They have a **student+teacher**
  - This provides the **self-distillation part**
  - **Momentum teacher -**
- The output of the **teacher** goes through **centering**, followed by a **softmax**
  - Keeps the model from **mode collapse**
- There's **no** use of **batch-norm**

### Self-Supervised Learning - you have no labels

- In this paper, *you don't have negative sample mechanism or contrastive learning mechanism*
  - **Negative Sampling Mechanism -**
  - **Contrastive Learning Mechanism -**
    - You take a few patches, A, a (say 2)
    - You then take a patch from another image b
    - **Patch A is your anchor**, and the model has to decide which **patch a or b** is in the image with **A**
    - **The model learns what kind of stuff is likely to be in the same image**
- **We want** a model that gives us **sensible representations**
  - Difficult with no labels
- We **augment images** in different ways
  - You do random perturbations to an image
  - You might flip it, add jitter, add some **solarisation** (dark parts lighter, light parts darker)

**Teachers only have global cropping**

**Students can have both**

- This means that whenever the student gets a small crop and teacher a large crop, the student has to learn that whatever small patch it sees is something part of a larger thing a teacher sees
  - Different to **contrastive learning**, you look at it **deeper** than just using **negative samples**
    - **You have to say what a small patch describes as a whole and differentiates it from others**

## Paper

- Utilises different augmentations
- You can also add **crops**
  - **Global Crops** - cover more than 50% of the image
  - **Local Crops** - less than 50% of the image

## Student-Teacher (**Self-Distillation**)

- We take an image and we **make two different augmentations of the same image**
- Make **two** different **versions** of it
- *We want the two models to output the same output for different augmentations of the same image*
- Our goal is to **minimise loss =  $-p_2 \log p_1$**
- We want the student and the teacher to **predict the same class output**
- *Theoretically, the easiest method would be to put them through the same transformer network*
  - **Can't pass the images through the same network because they would start predicting the same class**
    - This is called a **collapse** - when the model models everything to the **same representation**
- This idea comes from **distillation** - You have a **big model (teacher)** and you want to make the model **smaller**
  - You transfer the knowledge from the teacher model to the student model
  - **Works better than teaching the student from scratch**
- **Self-Distillation** - The two networks are of the same architecture
  - We **only train the student**
  - The **teacher** is constructed **from the student**
  - We train the student on the teacher until they predict the same thing, then **construct the teacher from the student** using **ema**
- **Exponentially moving average EMA** - we keep the teacher model, and as we update the student, **we move the teacher a little bit in the direction of the student model**
  - There is a scale associated with the step size
  - **Lots of hyperparameters**

---

### Algorithm 1 DINO PyTorch pseudocode w/o multi-crop.

---

```
# gs, gt: student and teacher networks
# C: center (K)
# tps, tpt: student and teacher temperatures
# l, m: network and center momentum rates
gt.params = gs.params
for x in loader: # load a minibatch x with n samples
    x1, x2 = augment(x), augment(x) # random views

    s1, s2 = gs(x1), gs(x2) # student output n-by-K
    t1, t2 = gt(x1), gt(x2) # teacher output n-by-K

    loss = H(t1, s2)/2 + H(t2, s1)/2
    loss.backward() # back-propagate

    # student, teacher and center updates
    update(gs) # SGD
    gt.params = l*gt.params + (1-l)*gs.params
    C = m*C + (1-m)*cat([t1, t2]).mean(dim=0)

def H(t, s):
    t = t.detach() # stop gradient
    s = softmax(s / tps, dim=1)
    t = softmax((t - C) / tpt, dim=1) # center + sharpen
    return - (t * log(s)).sum(dim=1).mean()
```

---

- **T1** = output of **x1** going through the **teacher**
- **S2** = output of **x2** going through the **student**

## Centering and Softmax

- You force the model to come up with a **k-dimensional** classification problem by itself
- *It chooses what the classes are so it has to make representations that **allows itself to come up with a problem that it can solve**, and centering + softmax makes sure it goes well and doesn't go into collapse*

**Centering** - The teacher keeps a running average of all of the representations that the teacher sees


- You **subtract** the average from the **logits**
  - Avoids **collapse**
  - Kind of like a **normalisation**
  - It keeps the **logits** be in a range that is manageable and have some variance
    - Since the student learns from the teacher, this also affects the student
- **Centering** prevents one dimension from dominating (always predicting the same class), however, this causes **collapse to a uniform distribution**
- Centering is only applied batch-wise and is like adding a **bias term** to the teacher

**Softmax (Sharpening)** -

- **Temperature Parameter** - the softmaxes in both networks have a temperature parameter
  - The two temperature parameters **are different**

$$P_s(x)^{(i)} = \frac{\exp(g_{\theta_s}(x)^{(i)} / \tau_s)}{\sum_{k=1}^K \exp(g_{\theta_s}(x)^{(k)} / \tau_s)},$$

- - The **temperature parameter**  $T_s$  for the **teacher** is much **lower**
    - This is called **sharpening**
- 
- 
- *Interesting they have softmax, couldn't they have just used l2-distance between logit representations*
- Here we use **cross entropy**
  - Softmax outputs a **normalised distribution**
  - You choose any number **k** to be the size of output
  - *The softmax will be a **distribution over the size of output** (where each node is a class)*
  - The teacher has **sharpening**, meaning it will have a more **peaked distribution**
    - Student gets a **less noisy, stronger signal**
- **Sharpening** has the opposite effect of **centering**
  - These two effects **balance their effects**



## Conclusion

- We **augment** the inputs in **two random ways**
- We put all of the images **through the student and the teacher**
- The teacher is an **ema of the student**
- This gives us **different representations of different augmentations** of the same image
- We put the **representations through a softmax** and force the output **distribution to be the same**
- The **teacher centers** the logits by an **exponential running average** of all of the representations it has seen
- The teacher also has a **sharper softmax**
- This results in **good representations and avoid collapse**

## Results

- We can now do **K-Nearest Neighbour classification**
  - A type of **zero-shot learning**
- **Image Retrieval**
- **Linear classification** on top of the representations
- **Copy Detection** - you wanna realise if someone made another image out of an original image
- **Ablation** - When comparing the attention heads for a ViT trained in a supervised way vs Dino, the Dino attention heads are way cleaner

The class representations come from the **CLS** token

- *You can visualise the attention heads of the **CLS token**, you get really strong semantic segmentation maps*
- A possible reason why **Dino attention is cleaner than supervised** is that:
  - Supervised models tend to stop learning when they master a problem
  - There is a problem of **shortcut learning**
  - The extra noise in attention maps is probably caused by **extra optimisations** that it tries to make that works overall for the whole dataset
  - There is no **hyper-optimisation** on a single task

## Why does this work?

- **Augmentations** are very important (i.e. the **multicrop augmentation** was very important)
  - **Multicrop Augmentation** -
  - The **augmentations** are essentially where you put the **human prior**
    - That's where you tell the model what to pay attention to and what to ignore
  - To do fully unsupervised, **domain agnostic** learning, we should replace **human designed augmentations**
  - Relies on **cropping** which is teaching the model to represent images **invariant to cropping position**
- **Dataset** - Datasets consist of **object-centered images** that always have an **object of interest**
  - **Strong Inductive Prior** - produces representations that focus on the objects of interests
  - *Makes sure that the **cropped images don't have different objects that have little in common***
  - As images are taken by humans, and you have to scrape the dataset somehow, the way you are creating a dataset is causing an **implicit bias** that affects where the **attention goes**
    - DINO probably relies a lot on **dataset construction**
    - We shouldn't expect this model to work on random pictures of the world, because the photos of the dataset used to train it **weren't IID**